Contact Information

No. 38 Tongyan Road, Jinnan District
Tianjin, P.R. China 300350

tnyang2000@gmail.com
https://tiannuo-yang.github.io/

Research Interests

As Moore's Law fades, what will be the next-generation AI system that can overcome the conflict between surging computational needs and scarce hardware resources? Inspired by this question, my current research focuses on **building low-cost, high-efficiency machine learning systems** that can be deployed to serve a wide range of scenarios.

Specifically, I identify performance bottlenecks, challenges, and opportunities in AI deployment, and leverage techniques such as online learning to address the performance- and resource-oriented problems. My works aim to harness hardware resources, automate the operation of complex systems, and enhance system efficiency (e.g., retrieval-augmented generation).

Education

University of Southern California, Los Angeles, CA, United States
Incoming Ph.D. Student in Computer Science from August 2025

Nankai University, Tianjin, China
Master in Computer Science (Exempted from Entrance Exam) August 2022 – Present
Advisor: Professor Yusen Li GPA: *3.63/4.0*
Thesis: Automated Performance Tuning Techniques for Parallel Applications

University of Science and Technology Beijing, Beijing, China
Bachelor in Information Management and Information System August 2018 – June 2022
Major GPA: *3.97/4.0* Cumulative GPA: *3.89/4.0*

Southern Taiwan University of Science and Technology, Taiwan, China
Major in Information Management (Exchange Program) September 2019 – January 2020
Cumulative GPA: *4.3/4.3*

Research Experience

University of Illinois at Urbana-Champaign, Urbana, IL, United States
Retrieval-Augmented Generation May 2024 – Present
Working with Professor Minjia Zhang on GPU-enhanced retrieval augmented generation, shedding lights on key concerns like batching strategies and latency-quality tradeoffs.

Nankai University, Tianjin, China
Datacenter, System, Machine Learning for System August 2022 – Present
Working with Professor Yusen Li on automatic performance tuning and hardware resource isolation for job collocations within multi-core systems.

Ant Group, Beijing, China
Vector Retrieval, Vector Database Optimization June 2023 – January 2024
Worked as a research intern under Dr. Jianguo Li and Wen Hu on optimizing AI infrastructure - vector database, enhancing CodeFuse services (a coding large language model platform).

University of Chinese Academy of Sciences, Beijing, China
Mixed Integer Programming, Heuristic Algorithm September 2020 – September 2021
Worked as an undergraduate research assistant under Professor Guanghui Zhou to develop data-driven combinatorial optimization problem (vehicle routing optimization).

Publications

Performance Tuning

[*] *Denotes equal contribution.* [†] *Denotes corresponding authorship.*

K. Cheng, Z. Wang, W. Hu, T. Yang, J. Li, S. Zhang. "SCOOT: Towards SLO-Optimized LLM Serving via Automatic Inference Engine Tuning." *The Web Conference (WWW)*, 2025. Oral Presentation.

T. Yang, W. Hu, W. Peng, Y. Li, J. Li, X. Liu, G. Wang. "VDTuner: Automated Performance Tuning for Vector Data Management Systems." *International Conference on Data Engineering (ICDE)*, 2024. Deployment on Ant Group's CodeFuse platform.

T. Yang, R. Chen, Y. Li, X. Liu, G. Wang. "CoTuner: A Hierarchical Learning Framework for Coordinately Optimizing Resource Partitioning and Parameter Tuning." *International Conference on Parallel Processing (ICPP)*, 2023.

Research Survey

Y. Zhou*, X. Lin*, X. Zhang*, M. Wang*, G. Jiang*, H. Lu*, Y. Wu*, K. Zhang*, Z. Yang*, K. Wang*, Y. Sui*, F. Jia* Z. Tang*, Y. Zhao*, H. Zhang*, T. Yang*, W. Chen*, Y. Mao*, Y. Li*, D. Bao*, Y. Li*, H. Liao*, T. Liu*, J. Liu*, J. Guo*, X. Zhao, Y. WEI, H. Qian, Q. Liu, X. Wang, W.K. Chan, C. Li, Y. Li, S. Yang, J. Yan, C. Mou, S. Han, W. Jin, G. Zhang, X. Zeng. "On the Opportunities of Green Computing: A Survey." *arXiv*, 2023
(Writing Section: 6.4 Resource Optimization).

Operations Research (Undergraduate Thesis)

T. Yang†, Z. Chu, B. Wang. "Feasibility on the Integration of Passenger and Freight Transportation in Rural Areas: A Service Mode and an Optimization Model." *Socio-Economic Planning Sciences* (SCI JCR Q1), 2023.

Research Projects

Qiyuan Laboratory Innovation Fund — November 2023 – November 2024
Worked as a core member on resource isolation mechanism for a multi-tenant cache system (i.e., Cachelib by Facebook).

CCF-Ant Research Fund on Green Computing — January 2023 – January 2024
Worked as a leader to write project proposal and conclusion, conduct research on AI infrastructure (vector database optimization) and practical platform deployment.

National Natural Science Foundation (NSF) of China — January 2023 – Present
Working as a core member on improving resource utilization in cloud with QoS guarantee.

Major Project of National NSF of China — December 2022 – Present
Working on real-time scheduling of cluster robots for major equipment manufacturing.

Honors And Awards

Ph.D. Student Fellowship, USC — 2025
Excellent Graduate, Nankai University (8 out of 157) — 2025
1st-Class Gongneng Scholarship, Nankai University — 2023, 2024
National 3rd Prize, Massive Storage Competition — 2022
National 2nd Prize (1493/41826 = top 3.6%), Contemporary Undergraduate Mathematical Contest in Modeling — 2020
Top Ten Singers on Campus, USTB — 2020

Talks And Services

Reviewer for
the 2025 ACM Web Conference.

Conference Talk at
the 40th ICDE at Utrecht, the Netherlands, May 2024;
the 52nd ICPP, Online, August 2023.

Teaching Assistant for

Computer Architecture (Fall 2023); C++ (Spring 2024).

Reading Group Founder and Leader for
Machine learning system research at Nankai-Baidu Joint Lab (from October 2024).

Open Source

SearchAgent-X: Highly efficient system for reasoning-search interleaved LLM agents.
https://github.com/tiannuo-yang/SearchAgent-X

VDTuner: Automated performance tuning framework for vector data management systems.
https://github.com/tiannuo-yang/VDTuner
(Chinese Blog: https://mp.weixin.qq.com/s/1JgXM5WSWBTv7fA0TLGfqw)

G-VRP-IPD-TW: Mathematical model and real-world dataset for a complex combinatorial optimization problem in transportation.
https://github.com/tiannuo-yang/G-VRP-IPD-TW